



Machine Learning


Exam 1 Review



Exam 1 Review

Classification

- ❑ Type of problem (binary, multi-class, multi-label)
- ❑ Cost functions, loss functions
- ❑ Metrics
- ❑ Regularization
- ❑ NN Output Activations
- ❑ Gradient Descent
 - ❑ Learning rate
 - ❑ Convergence vs Divergence



slido

Please download and install the
Slido app on all computers you use



A telecom company uses historical usage data to predict whether a specific customer will stay or cancel their subscription next month. What type of problem is this?

- ☐ Start presenting to display the poll results on this slide.



Which of the following best describes a loss function in machine learning?



A Junior Analyst, "Sam," is tasked with building a model to predict whether a specific customer will stay or cancel their subscription next month. Sam uses a Large Language Model (AI) to generate the code and presents the following model to you for approval before deploying it to production. Do you consider this to be a:



In the provided code, the use of `.toarray()` to convert a sparse matrix into a dense matrix will reduce the total amount of memory (RAM) consumed by the computer because it simplifies the data structure into a standard array format.



What is wrong with the configuration used by Sam?

slido

Please download and install the
Slido app on all computers you use



Which metric can you report for this problem?

- ☐ Start presenting to display the poll results on this slide.



Sam is training a neural network in TensorFlow for binary classification. He runs the code here. Is there anything wrong with the setup in the compile method?



Sam got this error. Based on your knowledge of Matrix Algebra and the error message above, why is the code failing?



slido

Please download and install the
Slido app on all computers you use



A churn prediction model outputs a probability of 0.85 that a customer will churn. The actual label is 1 (the customer churned). What is the binary cross-entropy loss for this single prediction?

☐ Start presenting to display the poll results on this slide.




You are given a model that predicts the probability of customer churn for four customers. The actual churn labels (1 = churn, 0 = no churn) and the model's predicted probabilities are shown in the table below. What is the binary cross-entropy cost for this batch?

- ☐ Start presenting to display the poll results on this slide.

Gradient Descent Algorithm

- Based on going in the direction of the decreasing gradient

$$\mathbf{w}_{i+1} = \mathbf{w}_i - \alpha \nabla J(\mathbf{w})$$



slido

Please download and install the
Slido app on all computers you use



**What is the most likely case
given the learning curve given
below**

☐ Start presenting to display the poll results on this slide.

slido

Please download and install the
Slido app on all computers you use



**What is the most likely case
given the learning curve below**

☐ Start presenting to display the poll results on this slide.



**What is the most likely case
given the learning curve of
the model fit by Sam**

- ☐ Start presenting to display the poll results on this slide.




What do you consider the main issue why this model is overfitting?



You are given a model that predicts the probability of customer churn for four customers. The actual churn labels (1 = churn, 0 = no churn). The cost function for binary crossentropy is 0.79. Suppose the model has the following three weights: $w = [0.8, -0.3, 1.2]$, and the regularization parameter $\alpha = 0.1$, what is the total cost: $J(w) + \alpha * ||w||_1$

Image Classification Example



Label	Description
0	T-shirt/top
1	Trouser
2	Pullover
3	Dress
4	Coat
5	Sandal
6	Shirt
7	Sneaker
8	Bag
9	Ankle boot

slido

Please download and install the
Slido app on all computers you use



What type of problem is this?

- ☐ Start presenting to display the poll results on this slide.

slido

Please download and install the
Slido app on all computers you use



What cost function(s) can we use?

☐ Start presenting to display the poll results on this slide.



slido

Please download and install the
Slido app on all computers you use



A Fashion MNIST model makes the following probability predictions for an image. The true class is “Sneaker” (Class 7).

☐ Start presenting to display the poll results on this slide.



slido

Please download and install the
Slido app on all computers you use



A Fashion MNIST model makes the following probability predictions for 4 images, with the actual class labels provided.

☐ Start presenting to display the poll results on this slide.

Types of Classification Problems

- The 3 main classification problems are:

Binary
Classification



- Spam
- Not spam

Multiclass
Classification



- Dog
- Cat
- Horse
- Fish
- Bird
- ...

Multi-label
Classification



- Dog
- Cat
- Horse
- Fish
- Bird
- ...



Which statement best describes multilabel classification?



In a neural network for multilabel classification with 5 possible labels, which output layer is most appropriate?



A multilabel dataset has 200 observations and 7 possible labels. What is the correct shape of the target matrix Y ?

Suppose the true labels for one observation are:

$[1, 0, 1, 0]$

and the predicted probabilities are:

$[0.91, 0.35, 0.62, 0.80]$

Using a threshold of 0.5, what is the predicted label vector?

