




Module 1

Regularization




The Gradient Descent Algorithm



Gradient Descent


- ❑ Gradient descent is an optimization algorithm used to find the minimum of a cost function.
- ❑ It is based on moving downwards the gradient of the cost function in steps.



$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta \mathbf{w}}$$

Gradient Descent

- ☐ Moving down the gradient.
- ☐ The negative gradient allows us to know how to update \mathbf{w}




$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta \mathbf{w}}$$

$$-\nabla J(\mathbf{w})$$

Gradient Descent

- ☐ Moving down the gradient.
- ☐ The negative gradient allows us to know how to update \mathbf{w}




$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta \mathbf{w}}$$

$$-\nabla J(\mathbf{w})$$

Gradient Descent

- If $\nabla J(\mathbf{w})$ is negative or positive we need to keep updating

$$-\nabla J(\mathbf{w})$$



If $\nabla J(\mathbf{w})$ is negative or positive we need to keep updating

Gradient Descent Algorithm


- ☐ The algorithm to update \mathbf{w}

$$\mathbf{w}_{i+1} = \mathbf{w}_i - \alpha \nabla J(\mathbf{w}_i)$$

Diagram illustrating the gradient descent update rule:

- \mathbf{w}_{i+1} : w at next step
- \mathbf{w}_i : w at current step
- step size: α
- gradient: $\nabla J(\mathbf{w}_i)$

The diagram shows a vertical arrow pointing upwards from \mathbf{w}_i to \mathbf{w}_{i+1} , representing the update. A horizontal arrow points from \mathbf{w}_i towards \mathbf{w}_{i+1} , representing the step size. A third horizontal arrow points from \mathbf{w}_i away from \mathbf{w}_{i+1} , representing the negative gradient.



Gradient Descent Algorithm


- ❑ The algorithm to update \mathbf{w}

$$\mathbf{w}_{i+1} = \mathbf{w}_i - \alpha \nabla J(\mathbf{w}_i)$$

Diagram illustrating the gradient descent update rule:

- \mathbf{w} at next step
- \mathbf{w} at current step
- step size
- gradient

The diagram shows the iterative update of parameters \mathbf{w} from the current step to the next step by subtracting the product of the step size and the gradient of the cost function J evaluated at the current parameters.



Gradient Descent Algorithm


- ❑ The algorithm to update \mathbf{w}

$$\mathbf{w}_{i+1} = \mathbf{w}_i - \alpha \nabla J(\mathbf{w}_i)$$

Diagram illustrating the gradient descent update rule:

- \mathbf{w} at next step
- \mathbf{w} at current step
- step size
- gradient

The diagram shows the formula $\mathbf{w}_{i+1} = \mathbf{w}_i - \alpha \nabla J(\mathbf{w}_i)$ where the gradient vector points downwards from the point \mathbf{w}_i towards the point \mathbf{w}_{i+1} .



Algorithm Main Steps

- ❑ Initialize \mathbf{w} with random values
- ❑ We take one step i at a time
- ❑ Measure the local gradient of the cost function $\nabla J(\mathbf{w}_i)$
- ❑ Update \mathbf{w} with direction of descent by the step size (learning rate)
- ❑ Stop when converging - the update is too small (say less than ϵ)
- ❑ Stop when diverging - maximum iterations reached (max_iter)


α : learning rate

$\nabla J(\mathbf{w}_i)$: gradient at value of \mathbf{w}_i

$\mathbf{w}_{i+1} = \mathbf{w}_i - \alpha \nabla J(\mathbf{w}_i)$: update




Elements of the GDA



Initializing w


- The parameters are initialized randomly
- Random normal initializations for weights are common
- The initialization affects the convergence to a minimum




Learning Rate - α

- This is the size of the step taken
- A too small learning rate can cause the algorithm to take too many steps to converge to the minimum
- A too large learning rate can overshoot (diverge from) the minimum

Convergence




Divergence




Gradients $\nabla J(\mathbf{w})$ Can be Zero at:

- A global minimum
- A local minimum
- A saddle point (neither a local or global minimum)



Number of Iterations:

- ❑ Gradient descent might continue to run infinitely if not stopped
- ❑ The maximum number of steps (**iterations**) are set in advance
- ❑ Most algorithms have parameters such as max_iter, iter, n_iter






Python

Example: Convex

- Suppose we want to estimate just one parameter w
- The cost function $J(w)$ is given below:


$$J(w) = w^2 - 4w$$



Example: Convex

- The gradient of the cost function $J(w)$ is given below:
- Let's also compute the gradient at $w = -1$

$$\nabla J(w) = 2w - 4$$
$$\nabla J(-1) = -6$$






Python

Example: Non-Convex

- ❑ Suppose we want to estimate just one parameter w
- ❑ The cost function $J(w)$ is given below:

$$J(w) = w^5 - 5w^4 + 5w^3 + 5w^2 - 6$$





Python

Gradient Descent for Regression

- Using MSE cost function:

$$\frac{1}{n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2$$


- Gradient

$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta \mathbf{w}} = \frac{2}{n} \mathbf{X}^\top (\mathbf{X}\mathbf{w} - \mathbf{y})$$

- GDA

$$\mathbf{w}_{n+1} = \mathbf{w}_n - \alpha \nabla J(\mathbf{w})$$

$$\mathbf{w}_{n+1} = \mathbf{w}_n - \alpha \frac{2}{n} \mathbf{X}^\top (\mathbf{X}\mathbf{w} - \mathbf{y})$$






Python

Optimizers:

- ❑ Get unstuck of local minima
- ❑ Speed up convergence



Batch Gradient Descent

- We can use parts of the data set to update w

Sales	Sq_Ft	lot
360000	3032	22221
340000	2058	22912
250000	1780	21345
205500	1638	17342
275500	2196	21786
248000	1966	18902
229900	2216	18639

y

X

$$\mathbf{w}_{n+1} = \mathbf{w}_n - \alpha \frac{2}{n} \mathbf{X}^\top (\mathbf{X}\mathbf{w} - \mathbf{y})$$

Batch Gradient Descent

- We can use parts of the data set to update w

Sales	Sq_Ft	lot
360000	3032	22221
340000	2058	22912
250000	1780	21345
205500	1638	17342
275500	2196	21786
248000	1966	18902
229900	2216	18639

$\{X_i, y_i\}_{i=1}^2$ - Batch = 2

$$\mathbf{w}_{n+1} = \mathbf{w}_n - \alpha \frac{2}{n} \mathbf{X}^\top (\mathbf{X}\mathbf{w} - \mathbf{y})$$

y

X

Batch Gradient Descent

- We can use parts of the data set to update w

Sales	Sq_Ft	lot
360000	3032	22221
340000	2058	22912
250000	1780	21345
205500	1638	17342
275500	2196	21786
248000	1966	18902
229900	2216	18639

y

X

$$\mathbf{w}_{n+1} = \mathbf{w}_n - \alpha \frac{2}{n} \mathbf{X}^\top (\mathbf{X}\mathbf{w} - \mathbf{y})$$

Batch Gradient Descent

- We can use parts of the data set to update w

Sales	Sq_Ft	lot
360000	3032	22221
340000	2058	22912
250000	1780	21345
205500	1638	17342
275500	2196	21786
248000	1966	18902
229900	2216	18639

y

X

$$\mathbf{w}_{n+1} = \mathbf{w}_n - \alpha \frac{2}{n} \mathbf{X}^\top (\mathbf{X}\mathbf{w} - \mathbf{y})$$

Batch Gradient Descent

- We can use parts of the data set to update w

Sales	Sq_Ft	lot
360000	3032	22221
340000	2058	22912
250000	1780	21345
205500	1638	17342
275500	2196	21786
248000	1966	18902
229900	2216	18639

y

X

$$\mathbf{w}_{n+1} = \mathbf{w}_n - \alpha \frac{2}{n} \mathbf{X}^\top (\mathbf{X}\mathbf{w} - \mathbf{y})$$

Completed 1 epoch

Stochastic Gradient Descent (SGD)

- We can use parts of the data set to update w

Sales	Sq_Ft	lot
360000	3032	22221
340000	2058	22912
250000	1780	21345
205500	1638	17342
275500	2196	21786
248000	1966	18902
229900	2216	18639

Mini batch of 1 observation


$$\mathbf{w}_{n+1} = \mathbf{w}_n - \alpha \frac{2}{n} \mathbf{X}^\top (\mathbf{X}\mathbf{w} - \mathbf{y})$$

y

X

Momentum


- ❑ Improves on gradient descent by using the acceleration of cost.
- ❑ Goal: get unstuck



	Derivative Form
Position	$r(t)$
Velocity	$v(t) = \frac{dr}{dt}$
Acceleration	$a(t) = \frac{dv}{dt} = \frac{d^2r}{dt^2}$

Adam Optimizer

- ❑ Uses both velocity and acceleration




```
Require:  $\alpha$ : Stepsize  
Require:  $\beta_1, \beta_2 \in [0, 1]$ : Exponential decay rates for the moment estimates  
Require:  $f(\theta)$ : Stochastic objective function with parameters  $\theta$   
Require:  $\theta_0$ : Initial parameter vector  
 $m_0 \leftarrow 0$  (Initialize 1st moment vector)  
 $v_0 \leftarrow 0$  (Initialize 2nd moment vector)  
 $t \leftarrow 0$  (Initialize timestep)  
while  $\theta_t$  not converged do  
     $t \leftarrow t + 1$   
     $g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$  (Get gradients w.r.t. stochastic objective at timestep  $t$ )  
     $m_t \leftarrow \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$  (Update biased first moment estimate)  
     $v_t \leftarrow \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$  (Update biased second raw moment estimate)  
     $\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$  (Compute bias-corrected first moment estimate)  
     $\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$  (Compute bias-corrected second raw moment estimate)  
     $\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$  (Update parameters)  
end while  
return  $\theta_t$  (Resulting parameters)
```



Python




Gradient Descent in LR




Example: Convex Cost Function

- ❑ Assume there is an unknown function $f: X \xrightarrow{\text{maps}} y$
- ❑ Assume a linear estimator (hypothesis) $h: X \xrightarrow{\text{maps}} y$ works well




Example: Convex Cost Function

- ❑ Assume there is a unknown function $f: \mathbf{X} \xrightarrow{\text{maps}} y$
- ❑ Assume a linear estimator (hypothesis) $h: \mathbf{X} \xrightarrow{\text{maps}} y$ works well
- ❑ If h is linear regression, then cost function SSE is convex



Example: Convex Cost Function

- ☐ We can find the parameter vector $\mathbf{w} = [w_0, w_1]$ minimizing SSE
- ☐ We do need to create a design matrix $\mathbf{X} = [\mathbf{1} \ \mathbf{x}_1]$
- ☐ The model predictions would be: $\mathbf{h}_{\mathbf{w}}(\mathbf{X}) = \mathbf{X} \mathbf{w}$



Example: Convex Cost Function

- ☐ We can find the parameter vector $\mathbf{w} = [w_0, w_1]$ minimizing SSE
- ☐ We do need to create a design matrix $\mathbf{X} = [\mathbf{1} \ \mathbf{x}_1]$
- ☐ The model predictions would be: $\mathbf{h}_{\mathbf{w}}(\mathbf{X}) = \mathbf{X}\mathbf{w}$

$$\begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ \vdots \\ \vdots \\ y_m \end{bmatrix} = \begin{bmatrix} 1 & x_{11} & x_{12} & \dots & \dots & x_{1n} \\ 1 & x_{21} & x_{22} & \dots & \dots & x_{2n} \\ 1 & x_{31} & x_{32} & \dots & \dots & x_{3n} \\ \vdots & \vdots & \vdots & \ddots & \ddots & \vdots \\ \vdots & \vdots & \vdots & \ddots & \ddots & \vdots \\ 1 & x_{m1} & x_{m2} & \dots & \dots & x_{mn} \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \\ \beta_3 \\ \vdots \\ \vdots \\ \beta_m \end{bmatrix} + \epsilon$$

1) Create y vector

2) Create X Matrix (also Called the Design/Model Matrix)

Example: Convex Cost Function


- We can find the parameter vector $\mathbf{w} = [w_0, w_1]$ minimizing SSE
- We do need to create a design matrix $\mathbf{X} = [\mathbf{1} \ \mathbf{x}_1]$
- The model predictions would be: $\mathbf{h}_{\mathbf{w}}(\mathbf{X}) = \mathbf{X} \ \mathbf{w}$

$$\mathbf{X} = \begin{bmatrix} 1 & X_{11} & X_{12} & \cdots & X_{1k} \\ 1 & X_{21} & X_{22} & \cdots & X_{2k} \\ \cdots & \cdots & \cdots & \cdots & \cdots \\ 1 & X_{n1} & X_{n2} & \cdots & X_{nk} \end{bmatrix} \quad \mathbf{w} = \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_k \end{bmatrix}$$

Example: Convex Cost Function

- Cost function is SSE


$$J(\mathbf{w}) = (\mathbf{y} - \mathbf{X}\mathbf{w})^\top(\mathbf{y} - \mathbf{X}\mathbf{w}) = \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2$$



Example: Convex Cost Function

- Gradient is the derivative of the cost function WRT parameters.

$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta \mathbf{w}} = -2(\mathbf{X}^\top \mathbf{y}) + 2(\mathbf{X}^\top \mathbf{X}\mathbf{w}) = 2\mathbf{X}^\top(\mathbf{X}\mathbf{w} - \mathbf{y})$$




Example: Convex Cost Function

- Setting the gradient to zero and solving for \mathbf{w} gets solution for min.

$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta \mathbf{w}} = -2(\mathbf{X}^\top \mathbf{y}) + 2(\mathbf{X}^\top \mathbf{X}\mathbf{w}) = 2\mathbf{X}^\top(\mathbf{X}\mathbf{w} - \mathbf{y})$$

Set to zero and solve for w
 $\mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1}(\mathbf{X}^\top \mathbf{y})$





Python