




Module 1

Regularization




The Gradient Descent Algorithm



Gradient Descent

- ❑ Gradient descent is an optimization algorithm used to find the minimum of a cost function.
- ❑ It is based on moving downwards the gradient of the cost function in steps.



$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta \mathbf{w}}$$

Gradient Descent

- ▣ Moving down the gradient.
- ▣ The negative gradient allows us to know how to update $\mathbf{w}$




$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta \mathbf{w}}$$

$$-\nabla J(\mathbf{w})$$

Gradient Descent

- ▣ Moving down the gradient.
- ▣ The negative gradient allows us to know how to update $\mathbf{w}$




$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta \mathbf{w}}$$

$$-\nabla J(\mathbf{w})$$

Gradient Descent

▣ If $\nabla J(\mathbf{w})$ is negative or positive we need to keep updating

$-\nabla J(\mathbf{w})$



If $\nabla J(\mathbf{w})$ is negative or positive we need to keep updating

Gradient Descent Algorithm

▣ The algorithm to update $\mathbf{w}$


$$\mathbf{w}_{i+1} = \mathbf{w}_i - \alpha \nabla J(\mathbf{w}_i)$$

$\mathbf{w}$ at next step

$\mathbf{w}$ at current step

step size

gradient



Gradient Descent Algorithm

▣ The algorithm to update $\mathbf{w}$


$$\mathbf{w}_{i+1} = \mathbf{w}_i - \alpha \nabla J(\mathbf{w}_i)$$

$\mathbf{w}$ at next step

$\mathbf{w}$ at current step

step size

gradient



Gradient Descent Algorithm

▣ The algorithm to update $\mathbf{w}$


$$\mathbf{w}_{i+1} = \mathbf{w}_i - \alpha \nabla J(\mathbf{w}_i)$$

$\mathbf{w}$ at next step

$\mathbf{w}$ at current step

step size

gradient



Algorithm Main Steps

- ▣ Initialize $\mathbf{w}$ with random values
- ▣ We take one step i at a time
- ▣ Measure the local gradient of the cost function $\nabla J(\mathbf{w}_i)$
- ▣ Update $\mathbf{w}$ with direction of descent by the step size (learning rate)
- ▣ Stop when converging - the update is too small (say less than ϵ)
- ▣ Stop when diverging - maximum iterations reached (max_iter)


α : learning rate

$\nabla J(\mathbf{w}_i)$: gradient at value of $\mathbf{w}_i$

$\mathbf{w}_{i+1} = \mathbf{w}_i - \alpha \nabla J(\mathbf{w}_i)$: update




Elements of the GDA



Initializing w


- ❑ The parameters are initialized randomly
- ❑ Random normal initializations for weights are common
- ❑ The initialization affects the convergence to a minimum




Learning Rate - α

- ❑ This is the size of the step taken
- ❑ A too small learning rate can cause the algorithm to take too many steps to converge to the minimum
- ❑ A too large learning rate can overshoot (diverge from) the minimum

Convergence




Divergence




Gradients $\nabla J(\mathbf{w})$ Can be Zero at:

- ☐ A global minimum
- ☐ A local minimum
- ☐ A saddle point (neither a local or global minimum)



Number of Iterations:

- ❑ Gradient descent might continue to run infinitely if not stopped
- ❑ The maximum number of steps (**iterations**) are set in advance
- ❑ Most algorithms have parameters such as `max_iter`, `iter`, `n_iter`






Python

Example: Convex

- ▣ Suppose we want to estimate just one parameter w
- ▣ The cost function $J(w)$ is given below:

$$J(w) = w^2 - 4w$$




Example: Convex

- ▣ The gradient of the cost function $J(w)$ is given below:
- ▣ Let's also compute the gradient at $w = -1$

$$\nabla J(w) = 2w - 4$$

$$\nabla J(-1) = -6$$






Python

Example: Non-Convex

- ▣ Suppose we want to estimate just one parameter w
- ▣ The cost function $J(w)$ is given below:

$$J(w) = w^5 - 5w^4 + 5w^3 + 5w^2 - 6$$





Python

Gradient Descent for Regression

- Using MSE cost function:

$$\frac{1}{n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2$$


- Gradient

$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta \mathbf{w}} = \frac{2}{n} \mathbf{X}^\top (\mathbf{X}\mathbf{w} - \mathbf{y})$$

- GDA

$$\mathbf{w}_{n+1} = \mathbf{w}_n - \alpha \nabla J(\mathbf{w})$$

$$\mathbf{w}_{n+1} = \mathbf{w}_n - \alpha \frac{2}{n} \mathbf{X}^\top (\mathbf{X}\mathbf{w} - \mathbf{y})$$






Python

Optimizers:

- ❑ Get unstuck of local minima
- ❑ Speed up convergence



Batch Gradient Descent

- ▣ We can use parts of the data set to update $\mathbf{w}$

Sales	Sq_Ft	lot
360000	3032	22221
340000	2058	22912
250000	1780	21345
205500	1638	17342
275500	2196	21786
248000	1966	18902
229900	2216	18639

y

X

$$\mathbf{w}_{n+1} = \mathbf{w}_n - \alpha \frac{2}{n} \mathbf{X}^\top (\mathbf{X} \mathbf{w} - \mathbf{y})$$

Batch Gradient Descent

- ▣ We can use parts of the data set to update $\mathbf{w}$

Sales	Sq_Ft	lot
360000	3032	22221
340000	2058	22912
250000	1780	21345
205500	1638	17342
275500	2196	21786
248000	1966	18902
229900	2216	18639

$\{X_i, y_i\}_{i=1}^2$ - Batch = 2

$$\mathbf{w}_{n+1} = \mathbf{w}_n - \alpha \frac{2}{n} \mathbf{X}^\top (\mathbf{X} \mathbf{w} - \mathbf{y})$$

$\mathbf{y}$

$\mathbf{X}$

Batch Gradient Descent

- ▣ We can use parts of the data set to update $\mathbf{w}$

Sales	Sq_Ft	lot
360000	3032	22221
340000	2058	22912
250000	1780	21345
205500	1638	17342
275500	2196	21786
248000	1966	18902
229900	2216	18639

y

X

$$\mathbf{w}_{n+1} = \mathbf{w}_n - \alpha \frac{2}{n} \mathbf{X}^\top (\mathbf{X} \mathbf{w} - \mathbf{y})$$

Batch Gradient Descent

- ▣ We can use parts of the data set to update $\mathbf{w}$

Sales	Sq_Ft	lot
360000	3032	22221
340000	2058	22912
250000	1780	21345
205500	1638	17342
275500	2196	21786
248000	1966	18902
229900	2216	18639

y

X

$$\mathbf{w}_{n+1} = \mathbf{w}_n - \alpha \frac{2}{n} \mathbf{X}^\top (\mathbf{X} \mathbf{w} - \mathbf{y})$$

Batch Gradient Descent

- ▣ We can use parts of the data set to update $\mathbf{w}$

Sales	Sq_Ft	lot
360000	3032	22221
340000	2058	22912
250000	1780	21345
205500	1638	17342
275500	2196	21786
248000	1966	18902
229900	2216	18639

y

X

$$\mathbf{w}_{n+1} = \mathbf{w}_n - \alpha \frac{2}{n} \mathbf{X}^\top (\mathbf{X} \mathbf{w} - \mathbf{y})$$

Completed 1 epoch

Stochastic Gradient Descent (SGD)

- ▣ We can use parts of the data set to update $\mathbf{w}$

Sales	Sq_Ft	lot
360000	3032	22221
340000	2058	22912
250000	1780	21345
205500	1638	17342
275500	2196	21786
248000	1966	18902
229900	2216	18639

Mini batch of 1 observation


$$\mathbf{w}_{n+1} = \mathbf{w}_n - \alpha \frac{2}{n} \mathbf{X}^\top (\mathbf{X} \mathbf{w} - y)$$

y

X

Momentum

- ❑ Improves on gradient descent by using the acceleration of cost.
- ❑ Goal: get unstuck



Derivative Form

Position

$$r(t)$$

Velocity


$$v(t) = \frac{dr}{dt}$$

Acceleration

$$a(t) = \frac{dv}{dt} = \frac{d^2r}{dt^2}$$

Adam Optimizer

❑ Uses both velocity and acceleration



Require: α : Stepsize

Require: $\beta_1, \beta_2 \in [0, 1]$: Exponential decay rates for the moment estimates

Require: $f(\theta)$: Stochastic objective function with parameters θ

Require: θ_0 : Initial parameter vector

$m_0 \leftarrow 0$ (Initialize 1st moment vector)

$v_0 \leftarrow 0$ (Initialize 2nd moment vector)

$t \leftarrow 0$ (Initialize timestep)

while θ_t not converged **do**

$t \leftarrow t + 1$

$g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$ (Get gradients w.r.t. stochastic objective at timestep t)

$m_t \leftarrow \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$ (Update biased first moment estimate)

$v_t \leftarrow \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$ (Update biased second raw moment estimate)

$\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$ (Compute bias-corrected first moment estimate)

$\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$ (Compute bias-corrected second raw moment estimate)

$\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$ (Update parameters)

end while


return θ_t (Resulting parameters)



Python




Gradient Descent in LR




Example: Convex Cost Function

- ❑ Assume there is an unknown function $f: \mathbf{X} \xrightarrow{\text{maps}} \mathbf{y}$
- ❑ Assume a linear estimator (hypothesis) $h: \mathbf{X} \xrightarrow{\text{maps}} \mathbf{y}$ works well




Example: Convex Cost Function

- ❑ Assume there is a unknown function $f: \mathbf{X} \xrightarrow{\text{maps}} \mathbf{y}$
- ❑ Assume a linear estimator (hypothesis) $h: \mathbf{X} \xrightarrow{\text{maps}} \mathbf{y}$ works well
- ❑ If h is linear regression, then cost function SSE is convex



Example: Convex Cost Function

- ▣ We can find the parameter vector $\mathbf{w} = [w_0, w_1]$ minimizing SSE
- ▣ We do need to create a design matrix $\mathbf{X} = [\mathbf{1} \ x_1]$
- ▣ The model predictions would be: $\mathbf{h}_{\mathbf{w}}(\mathbf{X}) = \mathbf{X} \mathbf{w}$



Example: Convex Cost Function

- ▣ We can find the parameter vector $\mathbf{w} = [w_0, w_1]$ minimizing SSE
- ▣ We do need to create a design matrix $\mathbf{X} = [\mathbf{1} \ \mathbf{x}_1]$
- ▣ The model predictions would be: $\mathbf{h}_{\mathbf{w}}(\mathbf{X}) = \mathbf{X} \mathbf{w}$

Bias or Intercept features

$$\begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ \vdots \\ \vdots \\ \vdots \\ y_m \end{bmatrix} = \begin{bmatrix} 1 & x_{11} & x_{12} & \dots & \dots & x_{1n} \\ 1 & x_{21} & x_{22} & \dots & \dots & x_{2n} \\ 1 & x_{31} & x_{32} & \dots & \dots & x_{3n} \\ \vdots & \vdots & \vdots & \dots & \dots & \vdots \\ \vdots & \vdots & \vdots & \dots & \dots & \vdots \\ \vdots & \vdots & \vdots & \dots & \dots & \vdots \\ 1 & x_{m1} & x_{m2} & \dots & \dots & x_{mn} \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \\ \beta_3 \\ \vdots \\ \vdots \\ \beta_m \end{bmatrix} + \epsilon$$

1) Create y vector

2) Create X Matrix (also Called the Design/Model Matrix)

Example: Convex Cost Function


- ▣ We can find the parameter vector $\mathbf{w} = [w_0, w_1]$ minimizing SSE
- ▣ We do need to create a design matrix $\mathbf{X} = [\mathbf{1} \ \mathbf{x}_1]$
- ▣ The model predictions would be: $\mathbf{h}_{\mathbf{w}}(\mathbf{X}) = \mathbf{X} \mathbf{w}$

$$\mathbf{X} = \begin{bmatrix} 1 & X_{11} & X_{12} & \cdots & X_{1k} \\ 1 & X_{21} & X_{22} & \cdots & X_{2k} \\ \cdots & \cdots & \cdots & \cdots & \cdots \\ 1 & X_{n1} & X_{n2} & \cdots & X_{nk} \end{bmatrix} \quad \mathbf{w} = \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_k \end{bmatrix}$$

Example: Convex Cost Function

- Cost function is SSE


$$J(\mathbf{w}) = (\mathbf{y} - \mathbf{X}\mathbf{w})^\top (\mathbf{y} - \mathbf{X}\mathbf{w}) = \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2$$



Example: Convex Cost Function

- Gradient is the derivative of the cost function WRT parameters.

$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta \mathbf{w}} = -2(\mathbf{X}^\top y) + 2(\mathbf{X}^\top \mathbf{X} \mathbf{w}) = 2\mathbf{X}^\top (\mathbf{X} \mathbf{w} - y)$$




Example: Convex Cost Function

- Setting the gradient to zero and solving for $\mathbf{w}$ gets solution for min.

$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta \mathbf{w}} = -2(\mathbf{X}^\top \mathbf{y}) + 2(\mathbf{X}^\top \mathbf{X} \mathbf{w}) = 2\mathbf{X}^\top (\mathbf{X} \mathbf{w} - \mathbf{y})$$

Set to zero and solve for $\mathbf{w}$

$$\mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1} (\mathbf{X}^\top \mathbf{y})$$





Python