




# Module 2

Binary Classification



# Binary Classification



# Types of Classification Problems

- The 3 main classification problems are:

Binary  
Classification



- Spam
- Not spam

Multiclass  
Classification



- Dog
- Cat
- Horse
- Fish
- Bird
- ...


Multi-label  
Classification



- Dog
- Cat
- Horse
- Fish
- Bird
- ...

# Binary Classification

- ❑ is the task of classifying the elements of a set into one of two groups (each called class).





# Binary Classification

- ❑ is the task of classifying the elements of a set into one of two groups (each called class).




# Binary Classification

- is the task of classifying the elements of a set into one of two groups (each called class).



# Binary Classification - Propensities

- Linear and non-linear predictors output the probability of belonging to a particular class ( $\hat{p}$ )



# Binary Cross-Entropy

- The most common **loss** function in binary classification

$$-(y \log \hat{p} + (1 - y) \log(1 - \hat{p}))$$

$y$

1
0
0
0

$\hat{p}$

0.775
0.116
0.039
0.070

How do we determine the **distance** between these two vectors?

# Binary Cross-Entropy

- The most common loss function in binary classification

$$-(y \log \hat{p} + (1 - y) \log(1 - \hat{p}))$$

$y$	$\hat{p}$	
1	0.775	$-(1 \log 0.775 + (1 - 1) \log(1 - 0.775)) = 0.2549$
0	0.116	
0	0.039	
0	0.070	

# Binary Cross-Entropy

- The most common loss function in binary classification

$$-(y \log \hat{p} + (1 - y) \log(1 - \hat{p}))$$

$y$	$\hat{p}$
1	0.775
0	0.116
0	0.039
0	0.070

$$-(0 \log 0.116 + (1 - 0) \log(1 - 0.116)) = 0.1233$$

# Binary Cross-Entropy

- The most common loss function in binary classification

$$-(y \log \hat{p} + (1 - y) \log(1 - \hat{p}))$$

$y$	$\hat{p}$
1	0.775
0	0.116
0	0.039
0	0.070

$$-(0 \log 0.039 + (1 - 0) \log(1 - 0.039)) = 0.0398$$

# Binary Cross-Entropy

- ❑ The most common loss function in binary classification

$$-(y \log \hat{p} + (1 - y) \log(1 - \hat{p}))$$

$y$	$\hat{p}$
1	0.775
0	0.116
0	0.039
0	0.070

$$-(0 \log 0.070 + (1 - 0) \log(1 - 0.070)) = 0.0726$$

# Binary Cross-Entropy

- ❑ The binary cross-entropy cost function is the average of the losses


$$J(w) = -\frac{1}{n} \sum (y \log \hat{p} + (1 - y) \log(1 - \hat{p}))$$

$y$	$\hat{p}$	
1	0.775	$0.2549 +$
0	0.116	$0.1233 +$
0	0.039	$0.0398 +$
0	0.070	$0.0726 = 0.4906/4 = 0.12265$

# Binary Cross-Entropy


- The binary cross-entropy cost function is the average of the losses

$$J(w) = -\frac{1}{n} \sum (y \log \hat{p} + (1 - y) \log(1 - \hat{p}))$$



Example: **LOGISTIC REGRESSION**

$$\hat{p} = \frac{1}{1 + e^{-(w_0 + w_1 x_1 + \dots)}} = \frac{1}{1 + e^{Xw}}$$





# Python

# Metrics

- Used to determine if model is performing well

Scoring Classification	Function	Comment
'accuracy'	<code>metrics.accuracy_score</code>	
'balanced_accuracy'	<code>metrics.balanced_accuracy_score</code>	
'top_k_accuracy'	<code>metrics.top_k_accuracy_score</code>	
'average_precision'	<code>metrics.average_precision_score</code>	
'neg_brier_score'	<code>metrics.brier_score_loss</code>	
'f1'	<code>metrics.f1_score</code>	for binary targets
'f1_micro'	<code>metrics.f1_score</code>	micro-averaged
'f1_macro'	<code>metrics.f1_score</code>	macro-averaged
'f1_weighted'	<code>metrics.f1_score</code>	weighted average
'f1_samples'	<code>metrics.f1_score</code>	by multilabel sample
'neg_log_loss'	<code>metrics.log_loss</code>	requires <code>predict_proba</code> support
'precision' etc.	<code>metrics.precision_score</code>	suffixes apply as with 'f1'
'recall' etc.	<code>metrics.recall_score</code>	suffixes apply as with 'f1'
'jaccard' etc.	<code>metrics.jaccard_score</code>	suffixes apply as with 'f1'
'roc_auc'	<code>metrics.roc_auc_score</code>	
'roc_auc_ovr'	<code>metrics.roc_auc_score</code>	
'roc_auc_ovo'	<code>metrics.roc_auc_score</code>	
'roc_auc_ovr_weighted'	<code>metrics.roc_auc_score</code>	
'roc_auc_ovo_weighted'	<code>metrics.roc_auc_score</code>	

# Metrics Review: Confusion Matrix

- ❑ The confusion matrix gives a summary of overall performance
- ❑ True Positives, True Negatives, False Positives and False Negatives

n=165	Predicted: NO	Predicted: YES	
Actual: NO	$n_{0,0}$ TN = 50	$n_{0,1}$ FP = 10	60
Actual: YES	$n_{1,0}$ FN = 5	$n_{1,1}$ TP = 100	105
	55	110	

# Metrics Review: Accuracy

- ❑ The accuracy metric measures number of obs predicted correctly
- ❑ It is one of the most important classification metrics

n=165	Predicted: NO	Predicted: YES	
Actual: NO	$n_{0,0}$ TN = 50	$n_{0,1}$ FP = 10	60
Actual: YES	$n_{1,0}$ FN = 5	$n_{1,1}$ TP = 100	105
	55	110	

$$\text{accuracy} = \frac{n_{1,1} + n_{0,0}}{n}$$

# Metrics Review: Recall (Sensitivity)

- ❑ The fraction of all 1's that are classified (predicted) as 1's.

n=165	Predicted: NO	Predicted: YES	
Actual: NO	$n_{0,0}$ TN = 50	$n_{0,1}$ FP = 10	60
Actual: YES	$n_{1,0}$ FN = 5	$n_{1,1}$ TP = 100	105
	55	110	

$$\text{sensitivity} = \frac{n_{1,1}}{n_{1,1} + n_{1,0}}$$

# Metrics Review: Precision

- The fraction of predicted values that are correct.

n=165	Predicted: NO	Predicted: YES	
Actual: NO	$n_{0,0}$ TN = 50	$n_{0,1}$ FP = 10	60
Actual: YES	$n_{1,0}$ FN = 5	$n_{1,1}$ TP = 100	105
	55	110	

$$\text{precision} = \frac{n_{1,1}}{n_{1,1} + n_{0,1}}$$

# F1 Score

- ❑ It is meant to capture both Recall and Precision

$$\text{F1 score} = 2 \frac{\text{precision} \times \text{recall}}{\text{precision} + \text{recall}}$$

n=165	Predicted: NO	Predicted: YES	
Actual: NO	$n_{0,0}$ TN = 50	$n_{0,1}$ FP = 10	60
Actual: YES	$n_{1,0}$ FN = 5	$n_{1,1}$ TP = 100	105
	55	110	

# Balanced Accuracy

- We use the balanced accuracy for imbalanced data

$$\text{balanced accuracy} = \frac{1}{2} \left( \frac{n_{0,0}}{n_{0,0} + n_{0,1}} + \frac{n_{1,1}}{n_{1,1} + n_{1,0}} \right)$$

n=165	Predicted: NO	Predicted: YES	
Actual: NO	$n_{0,0}$ TN = 50	$n_{0,1}$ FP = 10	60
Actual: YES	$n_{1,0}$ FN = 5	$n_{1,1}$ TP = 100	105
	55	110	



# Python


# Common Cost Functions Neural Nets

- ❑ Binary Cross-Entropy
  - ❑ Loss Function

$$J(w) = -\frac{1}{n} \sum (y \log \hat{p} + (1 - y) \log(1 - \hat{p}))$$

- ❑ Gradient

$$\nabla J(w) = \hat{p} - y$$





# Common Cost Functions Neural Nets

- ❑ Binary Focal Cross-Entropy
  - ❑ Loss Function

$$J(w) = -\frac{1}{n} \sum \alpha(1 - \hat{p})^\gamma y \log(\hat{p}) + (1 - \alpha)\hat{p}^\gamma(1 - y) \log(1 - \hat{p})$$


- ❑ Gradient

$$\nabla J(w) = (\gamma(1 - y - \hat{p}) \log(1 - \hat{p}) + (1 - y - \hat{p})\hat{p}(1 - \hat{p}))$$



# Activation Functions Neural Nets

- ❑ Common Activations on Layers
  - ❑ ReLu
  - ❑ Tanh
  - ❑ Sigmoid
- ❑ Output Layer Activation
  - ❑ Sigmoid
  - ❑ Tanh





# Python