




Module 2

Multi-Class Classification



Multi-Class Classification



Types of Classification Problems

- The 3 main classification problems are:

Binary
Classification



- Spam
- Not spam

Multiclass
Classification



- Dog
- Cat
- Horse
- Fish
- Bird
- ...


Multi-label
Classification



- Dog
- Cat
- Horse
- Fish
- Bird
- ...


Multi-Class Classification

- ❑ is a classification task with more than two classes.






Multi-Class Classification

- ❑ is a classification task with more than two classes.








Multi-Class Classification






Multi-Class Classification

Feature Matrix (X)





Multi-Class Classification




Multi-Class Classification

Feature Matrix (X)




Input ima



Multi-Class Classification - Propensities

- Linear and non-linear predictors output the probability of belonging to a particular class (\hat{p})



Categorical Cross-Entropy

- The most common **loss** function in multi-class classification

$$-\sum_k y_k \log \hat{p}_k$$

How do we determine the **distance** between these two matrices?

	DOG	CAT	BIRD
0	0	1	0
1	1	0	0
0	0	1	0
0	0	0	1

	p_dog	p_cat	p_bird
0.05	0.05	0.90	0.05
0.85	0.85	0.10	0.05
0.80	0.80	0.10	0.10
0.10	0.10	0.15	0.75

Categorical Cross-Entropy

- The most common loss function in multi-class classification

$$-\sum_k y_k \log \hat{p}_k$$

DOG	CAT	BIRD
0	1	0
1	0	0
0	1	0
0	0	1


p_dog	p_cat	p_bird
0.05	0.90	0.05
0.85	0.10	0.05
0.80	0.10	0.10
0.10	0.15	0.75

$$-(0 \log 0.05 + 1 \log 0.90 + 0 \log 0.05) = 0.1054$$

Categorical Cross-Entropy


- ❑ The cost function is averaged over the observations

$$J(w) = -\frac{1}{n} \sum_i \sum_k y_k \log \hat{p}_k$$



Combining Predictions

- We use the softmax function to combine predictions




$$\begin{bmatrix} 1.3 \\ 5.1 \\ 2.2 \\ 0.7 \\ 1.1 \end{bmatrix} \xrightarrow{\quad} \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}} \xrightarrow{\quad} \begin{bmatrix} 0.02 \\ 0.90 \\ 0.05 \\ 0.01 \\ 0.02 \end{bmatrix}$$

Binary Classification Algorithms?

- ❑ We can still use binary classification algorithms in multi-class classification
- ❑ Strategies: One-Vs-All (One-Vs-Rest) and One-Vs-One

Example: **LOGISTIC REGRESSION**


$$\hat{p} = \frac{1}{1 + e^{-(w_0 + w_1 x_1 + \dots)}} = \frac{1}{1 + e^{xw}}$$



One-Vs-All (One-Vs-Rest)



- This approach implies fitting the model multiple times
- In each fit the binary classifier compares one class against the rest

One-vs-all (one-vs-rest)



One-Vs-One

- This approach implies fitting the model multiple times
- In each fit the binary classifier compares one class against each other class



$$\hat{p} = \frac{1}{1 + e^{Xw}}$$

$$\hat{p} = \frac{1}{1 + e^{Xw}}$$

$$\hat{p} = \frac{1}{1 + e^{Xw}}$$

IRIS Dataset

- Classical ML dataset

iris setosa



petal sepal

iris versicolor



petal sepal


iris virginica



petal sepal

IRIS Dataset

- Classical ML dataset





Python


Metrics

- Used to determine if model is performing well

Scoring Classification	Function	Comment
'accuracy'	<code>metrics.accuracy_score</code>	
'balanced_accuracy'	<code>metrics.balanced_accuracy_score</code>	
'top_k_accuracy'	<code>metrics.top_k_accuracy_score</code>	
'average_precision'	<code>metrics.average_precision_score</code>	
'neg_brier_score'	<code>metrics.brier_score_loss</code>	
'f1'	<code>metrics.f1_score</code>	for binary targets
'f1_micro'	<code>metrics.f1_score</code>	micro-averaged
'f1_macro'	<code>metrics.f1_score</code>	macro-averaged
'f1_weighted'	<code>metrics.f1_score</code>	weighted average
'f1_samples'	<code>metrics.f1_score</code>	by multilabel sample
'neg_log_loss'	<code>metrics.log_loss</code>	requires <code>predict_proba</code> support
'precision' etc.	<code>metrics.precision_score</code>	suffixes apply as with 'f1'
'recall' etc.	<code>metrics.recall_score</code>	suffixes apply as with 'f1'
'jaccard' etc.	<code>metrics.jaccard_score</code>	suffixes apply as with 'f1'
'roc_auc'	<code>metrics.roc_auc_score</code>	
'roc_auc_ovr'	<code>metrics.roc_auc_score</code>	
'roc_auc_ovo'	<code>metrics.roc_auc_score</code>	
'roc_auc_ovr_weighted'	<code>metrics.roc_auc_score</code>	
'roc_auc_ovo_weighted'	<code>metrics.roc_auc_score</code>	

Metrics Review: Confusion Matrix


- ❑ The confusion matrix gives a summary of overall performance
- ❑ Accuracy can be derived as the fraction of correct predictions



$$\text{accuracy} = \frac{18 + 16 + 13}{50}$$

Metrics Review: Recall (Sensitivity)

- Detection rate for all classes
- We can use macro average (average by the number of classes).



$$\text{recall setosa} = \frac{18}{18} = 1$$


$$\text{recall versicolor} = \frac{16}{17} = 0.9412$$

$$\text{recall virginica} = \frac{13}{15} = 0.8667$$

$$\text{recall macro} = \frac{1 + 0.9412 + 0.8667}{3} = 0.9360$$

Metrics Review: Recall (Sensitivity)

- Detection rate for all classes
- We can use micro average (sum true detections over all obs).



$$\text{recall micro} = \frac{18 + 16 + 13}{18 + 17 + 15} = 0.9400$$



Python



Common Cost Functions Neural Nets

- ❑ Categorical Cross-Entropy
 - ❑ Loss Function

$$J(w) = - \sum_k y_k \log \hat{p}_k$$


- ❑ Gradient

$$\nabla J(w) = \hat{p} - y$$



Activation Functions Neural Nets

- ❑ Common Activations on Layers
 - ❑ ReLu
 - ❑ Tanh
 - ❑ Sigmoid
- ❑ Output Layer Activation
 - ❑ Softmax





Python