




Module 3

Deep Learning: Feed-Forward Architectures




DL: Feed-Forward Architectures



The Forward Pass


- ❑ Input nodes $\triangleright$ hidden nodes $\triangleright$ output nodes
- ❑ This is called the forward pass



$$\Sigma \longrightarrow w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4$$
$$f \longrightarrow f(w^T \mathbf{x})$$

The Forward Pass


- ❑ Input nodes $\triangleright$ hidden nodes $\triangleright$ output nodes
- ❑ This is called the forward pass




$$\begin{aligned}\Sigma &\longrightarrow b + w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 \\ &\qquad\qquad\qquad w^T x + b \\ f &\longrightarrow f(w^T x + b)\end{aligned}$$

The Forward Pass

- Input nodes $\triangleright$ hidden nodes $\triangleright$ output nodes
- This is called the forward pass



$$\begin{aligned}\Sigma &\longrightarrow 0 + .32x_1 + 0.3x_2 - 0.21x_3 + .002x_4 \\ &\quad \mathbf{w^T x + b} \\ f &\longrightarrow f(0 + .32x_1 + 0.3x_2 - 0.21x_3 + .002x_4)\end{aligned}$$



Parameter Initializers

- ❑ Setting initial values of model's weights and biases before training
 - ❑ Zero initializations
 - ❑ Advantages: computational complexity,
 - ❑ Disadvantages: all neurons learning similar features (symmetry learning)
 - ❑ Random Initialization
 - ❑ Normal
 - ❑ Uniform
 - ❑ Advantages: random initialization avoiding symmetry,
 - ❑ Disadvantages: if values are too large/small, issues with sigmoid and tanh activations
 - ❑ Xavier Glorot (tanh, sigmoid), He (ReLU), Lecun (sigmoid, tanh normalized)
 - ❑ Orthogonal (RNNs)

$$\Sigma \longrightarrow 0 + .32x_1 + 0.3x_2 - 0.21x_3 + .002x_4$$
$$w^T x + b$$

$$f \longrightarrow f(0 + .32x_1 + 0.3x_2 - 0.21x_3 + .002x_4)$$

Weights (kernels)

$$w = \begin{bmatrix} 0.320 \\ 0.300 \\ -0.210 \\ 0.002 \end{bmatrix}$$

Biases (intercepts)


$$b = [0]$$



Python


The Forward Pass

- ❑ Input nodes $\triangleright$ hidden nodes $\triangleright$ output nodes
- ❑ This is called the forward pass




The Forward Pass

- ❑ Input nodes $\triangleright$ hidden nodes $\triangleright$ output nodes
- ❑ This is called the forward pass



The Forward Pass

- Input nodes $\triangleright$ hidden nodes $\triangleright$ output nodes
- This is called the forward pass





	Bedrooms	Area	City_Distance	Age	Price
0	1	26.184098	1286.68	67	96004.804557
1	1	34.866901	1855.25	30	92473.722570
2	1	36.980709	692.09	24	98112.519940
3	1	17.445723	1399.49	66	92118.326874
4	1	52.587646	84.65	3	98976.653176



$$J(\mathbf{w}) = \frac{\sum_{i=1}^n (y - \hat{y})^2}{n}$$


The Forward Pass

- ❑ Input nodes $\triangleright$ hidden nodes $\triangleright$ output nodes
- ❑ This is called the forward pass




The Forward Pass

- Input nodes $\triangleright$ hidden nodes $\triangleright$ output nodes
- This is called the forward pass





$$\nabla J(\mathbf{w}) = 2\mathbf{X}^T(\mathbf{X}\mathbf{w} - \mathbf{y})$$





Backward Propagation

- ❑ Output nodes $\triangleright$ hidden nodes $\triangleright$ input nodes
- ❑ This is called the backward pass




Forward Pass

- ❑ Input nodes $\triangleright$ hidden nodes $\triangleright$ output nodes
- ❑ This is called the forward pass




Backward Pass

- ❑ Output nodes $\triangleright$ hidden nodes $\triangleright$ input nodes
- ❑ This is called the backward pass



$$\hat{y} = f(\mathbf{w}^T \mathbf{x} + b)$$





$$\mathbf{w}_{i+1} = \mathbf{w}_i - \alpha \nabla J(\mathbf{w})$$



Learning Curve

- ❑ Output nodes $\triangleright$ hidden nodes $\triangleright$ input nodes
- ❑ This is called the backward pass



slido

Please download and install the
Slido app on all computers you use



Consider the role of the learning rate in training a neural network. Which of the following statements best describes the effects of choosing a learning rate that is either too small or too large?

- ☐ Start presenting to display the poll results on this slide.

slido

Please download and install the
Slido app on all computers you use



Given a learning curve of the training set as given below, what is likely to be case in the training of this model.

- ☐ Start presenting to display the poll results on this slide.

slido

Please download and install the
Slido app on all computers you use



Given a learning curve of the training set as given below, what is likely to be case in the training of this model.


- ☐ Start presenting to display the poll results on this slide.




Python

Callbacks

- ❑ Special functions or methods that can be executed at specific points during the training
- ❑ Early stopping stops training when the validation loss stops improving



$$\mathbf{w}_{i+1} = \mathbf{w}_i - \alpha \nabla J(\mathbf{w})$$





Python

Optimizers (SGD vs Batch GD)

- ❑ We do not need to update the weights using ALL observations
- ❑ We still need to go over ALL observations (epoch)

Row ID	Study Hrs	Sleep Hrs	Quiz	Exam
1	12	6	78%	93%
2	22	6.5	24%	68%
3	115	4	100%	95%
4	31	9	67%	75%
5	0	10	58%	51%
6	5	8	78%	60%
7	92	6	82%	89%
8	57	8	91%	97%

Upd w's


Batch
Gradient
Descent

Row ID	Study Hrs	Sleep Hrs	Quiz	Exam
1	12	6	78%	93%
2	22	6.5	24%	68%
3	115	4	100%	95%
4	31	9	67%	75%
5	0	10	58%	51%
6	5	8	78%	60%
7	92	6	82%	89%
8	57	8	91%	97%

Upd w's

Stochastic
Gradient
Descent


$$\mathbf{w}_{i+1} = \mathbf{w}_i - \alpha \nabla J(\mathbf{w})$$




Optimizers (SGD)

- ❑ We do not need to update the weights using ALL observations
- ❑ We still need to go over ALL observations (epoch)

Stochastic Gradient Descent



$$\mathbf{w}_{i+1} = \mathbf{w}_i - \alpha \nabla J(\mathbf{w})$$



Batch Gradient Descent

- ❑ We do not need to update the weights using ALL observations
- ❑ A mini-batch can do as well
- ❑ We still need to go over ALL observations (epoch)

Row ID	Study Hrs	Sleep Hrs	Quiz	Exam
1	12	6	78%	93%
2	22	6.5	24%	68%
3	115	4	100%	95%
4	31	9	67%	75%
5	0	10	58%	51%
6	5	8	78%	60%
7	92	6	82%	89%
8	57	8	91%	97%

Upd w's


Batch
Gradient
Descent

Row ID	Study Hrs	Sleep Hrs	Quiz	Exam
1	12	6	78%	93%
2	22	6.5	24%	68%
3	115	4	100%	95%
4	31	9	67%	75%
5	0	10	58%	51%
6	5	8	78%	60%
7	92	6	82%	89%
8	57	8	91%	97%

Upd w's
Upd w's
Upd w's
Upd w's
Upd w's
Upd w's
Upd w's
Upd w's


Stochastic
Gradient
Descent

$$\mathbf{w}_{i+1} = \mathbf{w}_i - \alpha \nabla J(\mathbf{w})$$



Mini-Batch Gradient Descent

- ❑ We do not need to update the weights using ALL observations
- ❑ A mini-batch can do as well
- ❑ We still need to go over ALL observations (epoch)






Python

Forward Pass

- Hidden layers measure higher-order interactions




$$\hat{\mathbf{y}} = \mathbf{f}(\mathbf{w}^T \mathbf{x} + \mathbf{b})$$




Forward Pass

- Hidden layers measure higher-order interactions




$$\hat{y} = f(\mathbf{w}^T \mathbf{x} + \mathbf{b})$$




Update


- How do we update with activations and hidden layers?



$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta \mathbf{w}}$$




$$\hat{\mathbf{y}} = \mathbf{f}(\mathbf{w}^T \mathbf{x} + \mathbf{b})$$




Update


- How do we update with activations and hidden layers?



$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta f(z)} \frac{\delta f(z)}{\delta \mathbf{w}}$$




$$\hat{\mathbf{y}} = \mathbf{f}(\mathbf{w}^T \mathbf{x} + \mathbf{b})$$





Dropout Layers

- Regularization technique in neural networks that randomly sets a fraction of input units (neurons) to zero during training.




$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta f(\mathbf{z})} \frac{\delta f(\mathbf{z})}{\delta \mathbf{w}}$$

$$\hat{\mathbf{y}} = \mathbf{f}(\mathbf{w}^T \mathbf{x} + \mathbf{b})$$




L_1 (Lasso) / L_2 (Ridge) Layer Regularization


- Prevents overfitting by adding constraints to a model's parameters.



$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta f(\mathbf{z})} \frac{\delta f(\mathbf{z})}{\delta \mathbf{w}}$$



$$\hat{\mathbf{y}} = \mathbf{f}(\mathbf{w}^T \mathbf{x} + \mathbf{b})$$






Python

Everything is Vectorized


- ▣ Suppose $y = \{0,1\}$
- ▣ Cost function is the binary cross-entropy
- ▣ Output layer has sigmoid activation function




$$z = w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + b$$
$$a(z) = \frac{1}{1 + e^{-z}}$$

Everything is Vectorized

- ▣ Suppose $\mathbf{y} = \{0,1\}$
- ▣ Cost function is the binary cross-entropy
- ▣ Output layer has sigmoid activation function




$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta a(z)} \frac{\delta a(z)}{\delta \mathbf{w}}$$




Everything is Vectorized

- ▣ Suppose $\mathbf{y} = \{0,1\}$
- ▣ Cost function is the binary cross-entropy
- ▣ Output layer has sigmoid activation function




$$\frac{\delta J(\mathbf{w})}{\delta a(z)} = \frac{-\delta \sum_{i=1}^n (y_i \log(a) + (1 - y_i) \log(1 - a))}{\delta a}$$

$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta a(z)} \frac{\delta a(z)}{\delta \mathbf{w}}$$




Everything is Vectorized

- ▣ Suppose $y = \{0,1\}$
- ▣ Cost function is the binary cross-entropy
- ▣ Output layer has sigmoid activation function




$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta a(z)} \frac{\delta a(z)}{\delta \mathbf{w}}$$




Everything is Vectorized

- ▣ Suppose $\mathbf{y} = \{0,1\}$
- ▣ Cost function is the binary cross-entropy
- ▣ Output layer has sigmoid activation function




$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta a(z)} \frac{\delta a(z)}{\delta \mathbf{w}}$$




Everything is Vectorized

- ▣ Suppose $\mathbf{y} = \{0,1\}$
- ▣ Cost function is the binary cross-entropy
- ▣ Output layer has sigmoid activation function




$$\frac{\delta J(\mathbf{w})}{\delta \mathbf{w}} = (a - y)\mathbf{x}^\top$$

$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta a(z)} \frac{\delta a(z)}{\delta \mathbf{w}}$$




Everything is Vectorized

- ▣ Suppose $\mathbf{y} = \{0,1\}$
- ▣ Cost function is the binary cross-entropy
- ▣ Output layer has sigmoid activation function




$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta a(z)} \frac{\delta a(z)}{\delta \mathbf{w}}$$



Everything is Vectorized


- ▣ Suppose $\mathbf{y} = \{0,1\}$
- ▣ Cost function is the binary cross-entropy
- ▣ Output layer has sigmoid activation function



How weights are updated


$$\nabla J(\mathbf{w}) = (\mathbf{A} - \mathbf{Y})\mathbf{X}^T$$

$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta a(z)} \frac{\delta a(z)}{\delta \mathbf{w}}$$



Everything is Vectorized

- ▣ Suppose $\mathbf{y} = \{0,1\}$
- ▣ Cost function is the binary cross-entropy
- ▣ Output layer has sigmoid activation function




How biases are updated

$$\nabla J(\mathbf{w}) = (A - Y)\mathbf{1}^\top$$

All updates are done here

$$\mathbf{w}_{i+1} = \mathbf{w}_i - \alpha \nabla J(\mathbf{w})$$

$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta a(z)} \frac{\delta a(z)}{\delta \mathbf{w}}$$





Python