

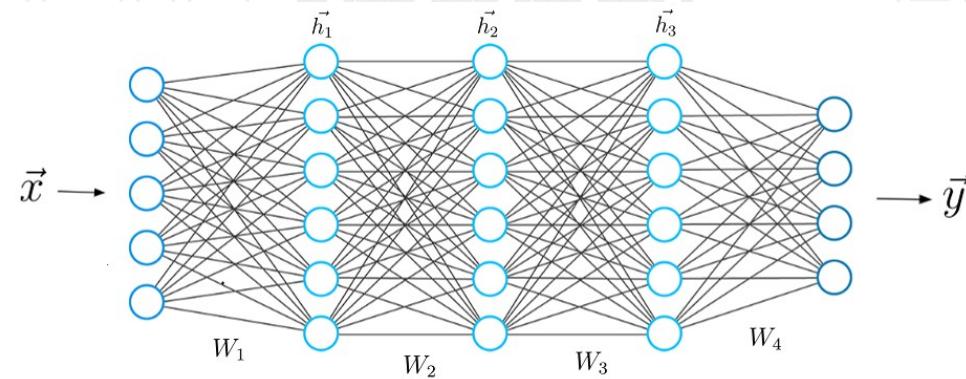


# Module 3

Deep Learning: Feed-Forward Architectures




# DL: Feed-Forward Architectures




# The Forward Pass

- Input nodes  $\triangleright$  hidden nodes  $\triangleright$  output nodes
- This is called the forward pass




# The Forward Pass

- Input nodes  $\triangleright$  hidden nodes  $\triangleright$  output nodes
- This is called the forward pass




# The Forward Pass

- Input nodes ▷ hidden nodes ▷ output nodes
- This is called the forward pass



$$\begin{aligned}\Sigma &\longrightarrow 0 + .32x_1 + 0.3x_2 - 0.21x_3 + .002x_4 \\ w^\top x + b \\ f &\longrightarrow f(0 + .32x_1 + 0.3x_2 - 0.21x_3 + .002x_4)\end{aligned}$$



# Parameter Initializers

- ❑ Setting initial values of model's weights and biases before training
  - ❑ Zero initializations
    - ❑ Advantages: computational complexity,
    - ❑ Disadvantages: all neurons learning similar features (symmetry learning)
  - ❑ Random Initialization
    - ❑ Normal
    - ❑ Uniform
    - ❑ Advantages: random initialization avoiding symmetry,
    - ❑ Disadvantages: if values are too large/small, issues with sigmoid and tanh activations
  - ❑ Xavier Glorot (tanh, sigmoid), He (ReLU), Lecun (sigmoid, tanh normalized)
  - ❑ Orthogonal (RNNs)

$$\begin{aligned}\Sigma \longrightarrow & 0 + .32x_1 + 0.3x_2 - 0.21x_3 + .002x_4 \\ & \mathbf{w}^T \mathbf{x} + \mathbf{b} \\ f \longrightarrow & f(0 + .32x_1 + 0.3x_2 - 0.21x_3 + .002x_4)\end{aligned}$$


Weights (kernels)	Biases (intercepts)
$\mathbf{w} = \begin{bmatrix} 0.320 \\ 0.300 \\ -0.210 \\ 0.002 \end{bmatrix}$	$\mathbf{b} = [0]$



# Python


# The Forward Pass

- Input nodes  $\triangleright$  hidden nodes  $\triangleright$  output nodes
- This is called the forward pass




# The Forward Pass

- Input nodes  $\triangleright$  hidden nodes  $\triangleright$  output nodes
- This is called the forward pass




# The Forward Pass

- Input nodes  $\triangleright$  hidden nodes  $\triangleright$  output nodes
- This is called the forward pass






Bedrooms	Area	City_Distance	Age	Price
0	1	26.184098	1286.68	96004.804557
1	1	34.866901	1855.25	92473.722570
2	1	36.980709	692.09	98112.519940
3	1	17.445723	1399.49	92118.326874
4	1	52.587646	84.65	98976.653176



$$J(w) = \frac{\sum_{i=1}^n (y - \hat{y})^2}{n}$$


# The Forward Pass

- Input nodes  $\triangleright$  hidden nodes  $\triangleright$  output nodes
- This is called the forward pass





# The Forward Pass

- Input nodes  $\triangleright$  hidden nodes  $\triangleright$  output nodes
- This is called the forward pass






$$\nabla J(\mathbf{w}) = 2\mathbf{X}^T(\mathbf{X}\mathbf{w} - \mathbf{y})$$






# Backward Propagation

- Output nodes  $\triangleright$  hidden nodes  $\triangleright$  input nodes
- This is called the backward pass






# Forward Pass

- Input nodes  $\triangleright$  hidden nodes  $\triangleright$  output nodes
- This is called the forward pass





# Backward Pass

- Output nodes  $\triangleright$  hidden nodes  $\triangleright$  input nodes
- This is called the backward pass



# Learning Curve

- ❑ Output nodes  $\triangleright$  hidden nodes  $\triangleright$  input nodes
- ❑ This is called the backward pass



Please download and install the Slido app on all computers you use



**Consider the role of the learning rate in training a neural network. Which of the following statements best describes the effects of choosing a learning rate that is either too small or too large?**

- Start presenting to display the poll results on this slide.

Please download and install the Slido app on all computers you use



**Given a learning curve of the training set as given below, what is likely to be case in the training of this model.**

- Start presenting to display the poll results on this slide.

Please download and install the Slido app on all computers you use



**Given a learning curve of the training set as given below, what is likely to be case in the training of this model.**


- Start presenting to display the poll results on this slide.




# Python

# Callbacks

- ❑ Special functions or methods that can be executed at specific points during the training
- ❑ Early stopping stops training when the validation loss stops improving



$$\mathbf{w}_{i+1} = \mathbf{w}_i - \alpha \nabla J(\mathbf{w})$$





# Python

# Optimizers (SGD vs Batch GD)

- ❑ We do not need to update the weights using ALL observations
- ❑ We still need to go over ALL observations (epoch)

Upd w's ←

Row ID	Study Hrs	Sleep Hrs	Ouiz	Exam
1	12	6	78%	93%
2	22	6.5	24%	68%
3	115	4	100%	95%
4	31	9	67%	75%
5	0	10	58%	51%
6	5	8	78%	60%
7	92	6	82%	89%
8	57	8	91%	97%


Batch  
Gradient  
Descent

Upd w's ←

Row ID	Study Hrs	Sleep Hrs	Quiz	Exam
1	12	6	78%	93%
2	22	6.5	24%	68%
3	115	4	100%	95%
4	31	9	67%	75%
5	0	10	58%	51%
6	5	8	78%	60%
7	92	6	82%	89%
8	57	8	91%	97%

Stochastic  
Gradient  
Descent


$$\mathbf{w}_{i+1} = \mathbf{w}_i - \alpha \nabla J(\mathbf{w})$$




# Optimizers (SGD)

- We do not need to update the weights using ALL observations
- We still need to go over ALL observations (epoch)

**Stochastic** Gradient  
Descent



$$\mathbf{w}_{i+1} = \mathbf{w}_i - \alpha \nabla J(\mathbf{w})$$



# Batch Gradient Descent

- ❑ We do not need to update the weights using ALL observations
- ❑ A mini-batch can do as well
- ❑ We still need to go over ALL observations (epoch)

Upd w's ←

Row ID	Study Hrs	Sleep Hrs	Ouiz	Exam
1	12	6	78%	93%
2	22	6.5	24%	68%
3	115	4	100%	95%
4	31	9	67%	75%
5	0	10	58%	51%
6	5	8	78%	60%
7	92	6	82%	89%
8	57	8	91%	97%


  

Row ID	Study Hrs	Sleep Hrs	Quiz	Exam
1	12	6	78%	93%
2	22	6.5	24%	68%
3	115	4	100%	95%
4	31	9	67%	75%
5	0	10	58%	51%
6	5	8	78%	60%
7	92	6	82%	89%
8	57	8	91%	97%

Batch  
Gradient  
Descent


Stochastic  
Gradient  
Descent

$$\mathbf{w}_{i+1} = \mathbf{w}_i - \alpha \nabla J(\mathbf{w})$$



# Mini-Batch Gradient Descent

- We do not need to update the weights using ALL observations
- A mini-batch can do as well
- We still need to go over ALL observations (epoch)







# Python

# Forward Pass

- Hidden layers measure higher-order interactions




$$\hat{y} = f(\mathbf{w}^T \mathbf{x} + b)$$





# Forward Pass

- Hidden layers measure higher-order interactions




$$\hat{y} = f(\mathbf{w}^T \mathbf{x} + b)$$





# Update

- How do we update with activations and hidden layers?




$$\nabla J(w) = \frac{\delta J(w)}{\delta w}$$
$$\hat{y} = f(w^T x + b)$$




# Update


- How do we update with activations and hidden layers?



$$\nabla J(w) = \frac{\delta J(w)}{\delta f(z)} \frac{\delta f(z)}{\delta w}$$




$$\hat{y} = f(w^T x + b)$$





# Dropout Layers

- Regularization technique in neural networks that randomly sets a fraction of input units (neurons) to zero during training.




$$\nabla J(w) = \frac{\delta J(w)}{\delta f(z)} \frac{\delta f(z)}{\delta w}$$

$$\hat{y} = f(w^T x + b)$$





# $L_1$ (Lasso) / $L_2$ (Ridge) Layer Regularization

- Prevents overfitting by adding constraints to a model's parameters.



$$\nabla J(w) = \frac{\delta J(w)}{\delta f(z)} \frac{\delta f(z)}{\delta w}$$

$$\hat{y} = f(w^T x + b)$$






# Python

# Everything is Vectorized


- Suppose  $y = \{0,1\}$
- Cost function is the binary cross-entropy
- Output layer has sigmoid activation function




$$z = w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + b$$
$$a(z) = \frac{1}{1+e^{-z}}$$

# Everything is Vectorized

- Suppose  $y = \{0,1\}$
- Cost function is the binary cross-entropy
- Output layer has sigmoid activation function




$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta a(z)} \frac{\delta a(z)}{\delta \mathbf{w}}$$




# Everything is Vectorized

- Suppose  $y = \{0,1\}$
- Cost function is the binary cross-entropy
- Output layer has sigmoid activation function




$$\frac{\delta J(\mathbf{w})}{\delta a(z)} = \frac{-\delta \sum_{i=1}^n (y_i \log(a) + (1 - y_i)\log(1 - a))}{\delta a}$$

$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta a(z)} \frac{\delta a(z)}{\delta \mathbf{w}}$$




# Everything is Vectorized

- Suppose  $y = \{0,1\}$
- Cost function is the binary cross-entropy
- Output layer has sigmoid activation function




$$\frac{\delta J(\mathbf{w})}{\delta a(z)} = a - y$$

$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta a(z)} \frac{\delta a(z)}{\delta \mathbf{w}}$$




# Everything is Vectorized

- Suppose  $y = \{0,1\}$
- Cost function is the binary cross-entropy
- Output layer has sigmoid activation function




$$\frac{\delta a(z)}{\delta \mathbf{w}} = \frac{\delta \mathbf{w}^\top \mathbf{x} + b}{\delta \mathbf{w}} = \mathbf{x}^\top$$

$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta a(z)} \frac{\delta a(z)}{\delta \mathbf{w}}$$




# Everything is Vectorized

- Suppose  $y = \{0,1\}$
- Cost function is the binary cross-entropy
- Output layer has sigmoid activation function




$$\frac{\delta J(\mathbf{w})}{\delta \mathbf{w}} = (a - y)\mathbf{x}^\top$$

$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta a(z)} \frac{\delta a(z)}{\delta \mathbf{w}}$$




# Everything is Vectorized

- Suppose  $y = \{0,1\}$
- Cost function is the binary cross-entropy
- Output layer has sigmoid activation function




$$\nabla J(\mathbf{w}) = (a - y)\mathbf{x}^\top$$

$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta a(z)} \frac{\delta a(z)}{\delta \mathbf{w}}$$




# Everything is Vectorized

- Suppose  $y = \{0,1\}$
- Cost function is the binary cross-entropy
- Output layer has sigmoid activation function




How weights are updated  
 $\nabla J(\mathbf{w}) = (A - Y)\mathbf{X}^\top$

$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta a(z)} \frac{\delta a(z)}{\delta \mathbf{w}}$$



# Everything is Vectorized

- Suppose  $y = \{0,1\}$
- Cost function is the binary cross-entropy
- Output layer has sigmoid activation function




How biases are updated

$$\nabla J(\mathbf{w}) = (A - Y)\mathbf{1}^\top$$

All updates are done here

$$\mathbf{w}_{i+1} = \mathbf{w}_i - \alpha \nabla J(\mathbf{w})$$

$$\nabla J(\mathbf{w}) = \frac{\delta J(\mathbf{w})}{\delta a(z)} \frac{\delta a(z)}{\delta \mathbf{w}}$$





# Python