

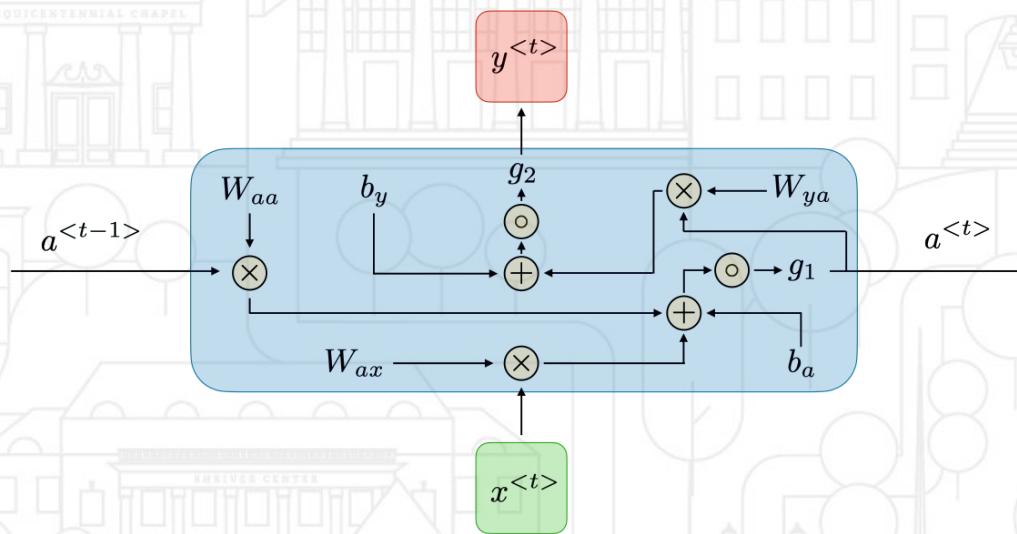


Module 6

Recurrent Neural Networks



DL: RNNs




Sequences

- ❑ Enumerated collection of objects
- ❑ Repetitions in sequences are allowed and order matters.
- ❑ Time series, speech rely on changes over a time interval

$$\{a_n\} = \{a_1, a_2, a_3, \dots, a_n\}$$


Sequences Examples

- Stock price, financial instruments, inventory predictions





Sequences Examples

- ❑ Natural language processing, LLMs



Univariate vs Multivariate Sequences

- We may need to understand more than one sequence



Feature Creation

- We can use a window size (width, length, step) and horizon
- This method is called the “sliding window”



```
# Window for one week with the target of predicting the next day (Bitcoin prices)
[123.654, 125.455, 108.584, 118.674, 121.338, 120.655, 121.795] -> [123.033]
[125.455, 108.584, 118.674, 121.338, 120.655, 121.795, 123.033] -> [124.049]
[108.584, 118.674, 121.338, 120.655, 121.795, 123.033, 124.049] -> [125.961]
```

window size = 7, horizon =
1

Feature Creation

- Using a window size (input width) and horizon (label width)
- This method is called the sliding window

```
time, measure
1, 100
2, 110
3, 108
4, 115
5, 120
```

LAG 1 feature (window size = 1, horizon = 1)

```
X, y
?, 100
100, 110
110, 108
108, 115
115, 120
120, ?
```

Feature Creation

- Using a window size (input width) and horizon (label width)
- This method is called the sliding window

```
time, measure
1, 100
2, 110
3, 108
4, 115
5, 120
```

LAG 1 feature (window size = 1, horizon = 2)

```
X1, y1, y2
?, 100, 110
100, 110, 108
110, 108, 115
108, 115, 120
115, 120, ?
120, ?, ?
```



Python

Creating Lags on Multiple Features

- Using the sliding window method on multiple predictors can help
- The windows used can be different

```
time, measure
1, 100
2, 110
3, 108
4, 115
5, 120
```

LAG 1 feature (window size = 1, horizon = 2)


```
X1, y1, y2
?, 100, 110
100, 110, 108
110, 108, 115
108, 115, 120
115, 120, ?
120, ?, ?
```




Python

Recurrent Neural Networks

- ❑ RNN neurons contain a feedback loop to receive sequence inputs
- ❑ The neurons (units) can be rolled to account for a sequence.




$$y = f(w_{ih}x + b)$$



$$h_t = f(w_{ih}x + w_{hh}h_{t-1} + b)$$

Recurrent Neural Networks


- ❑ RNN neurons contain a feedback loop to receive sequence inputs
- ❑ The neurons (units) can be rolled to account for a sequence.



$$h_t = f(w_{ih}x + w_{hh}h_{t-1} + b)$$


Recurrent Neural Networks

- ❑ RNN neurons contain a feedback loop to receive sequence inputs
- ❑ The neurons (units) can be rolled to account for a sequence.




Recurrent Neural Networks

- Sequence to Sequence (seq2seq)




Recurrent Neural Networks

- Sequence to Vector (seq2vec)




Recurrent Neural Networks

- Vector to Sequence (vec2seq)




Recurrent Neural Networks

- ❑ Elman network
- ❑ Tanh activation




Stacked RNN Architecture

- Stacked (deep) RNNs are composed of multiple RNNs stacked one above the other.



RNN Architecture

- We can construct layers of RNN neurons (units)
- We can also stack layers (need to return the sequences)




General Form of RNNs



Python

Bi-directional RNN Architecture

- ❑ Composed of two RNNs
- ❑ The input sequence (RNN # 1) and opposite direction (RNN # 2)






Python

CNN for Sequences

- ❑ CNNs excel at capturing **local dependencies** (e.g., trends, peaks).
- ❑ Effective in **repeating patterns** regardless of position in sequence.
- ❑ Fewer parameters than RNNs; allows **parallel computation**.



TimeSeriesGenerator

- ❑ Converts raw time series into supervised learning **format**
- ❑ Efficient for memory use and model training
- ❑ Supports **sliding window** logic

python

 Copy

```
TimeseriesGenerator(  
    data,                      # Numpy array of time series  
    targets,                   # Numpy array of target values  
    length=10,                 # Number of time steps per input sample  
    sampling_rate=1,            # Period between samples  
    stride=1,                  # Step between successive windows  
    batch_size=32,              # Number of samples per batch  
    shuffle=False               # Whether to shuffle samples  
)
```

TimeSeriesGenerator

- Converts raw time series into supervised learning format
- Efficient for memory use and model training
- Supports **sliding window** logic

Time	Stock	Feature 1	Feature 2	Feature ...	Feature N			
T+0	1	[some value]	[some value]	[some value]	[some value]	Sequence 1	Sequence 2	Sequence 3
T+1		[some value]	[some value]	[some value]	[some value]			
T+2		[some value]	[some value]	[some value]	[some value]			
T+3		[some value]	[some value]	[some value]	[some value]			
T+4		[some value]	[some value]	[some value]	[some value]			
T+5	1	[some value]	[some value]	[some value]	[some value]			
T+0	2	[some value]	[some value]	[some value]	[some value]			
T+1	2	[some value]	[some value]	[some value]	[some value]			
T+2	2	[some value]	[some value]	[some value]	[some value]			
T+3	2	[some value]	[some value]	[some value]	[some value]			
T+4	2	[some value]	[some value]	[some value]	[some value]			
T+5	2	[some value]	[some value]	[some value]	[some value]			
T+0	3	[some value]	[some value]	[some value]	[some value]			
T+1	3	[some value]	[some value]	[some value]	[some value]			
T+2	3	[some value]	[some value]	[some value]	[some value]			

Review: Convolutions

- Suppose we have a sequence $I_{n \times p}$ and a kernel $K_{k \times l}$
- The resulting image $O_{n-k+1} \times p-l+1$

7	2	3	3	8
4	5	3	8	4
3	3	2	8	4
2	8	7	2	7
5	4	4	5	4

*

1	0	-1
1	0	-1
1	0	-1

=

6		

$$\begin{aligned} & 7 \times 1 + 4 \times 1 + 3 \times 1 + \\ & 2 \times 0 + 5 \times 0 + 3 \times 0 + \\ & 3 \times -1 + 3 \times -1 + 2 \times -1 \\ & = 6 \end{aligned}$$

$I_{5 \times 5}$

$K_{3 \times 3}$

$O_{5-3+1 \times 5-3+1} = O_{3 \times 3}$

Review: Convolutions

- ❑ Suppose we have a sequence $I_{n \times p}$ and a kernel $K_{k \times l}$
- ❑ The resulting image $O_{n-k+1} \times p-l+1$



Review: Pooling

- ❑ Pooling is also used to downsample the sequences.
- ❑ Pooling filters keep the important parts of the sequence.
- ❑ Max, Min and Average Pooling Filters are the most common

0	0	1	2	2	1	0
0	0	2	2	3	1	0
1	1	3	4	3	2	1
0	1	3	1	3	1	0
1	2	2	5	2	2	1
0	1	2	2	2	1	0
0	1	1	2	1	1	0

Max
→

0		



Python

Limitations of RNNs

- ❑ Struggle with long-term dependencies due to vanishing gradients.
 - ❑ Attention (weights or scores for each hidden state)
- ❑ Training can be slow and unstable.
- ❑ Can be replaced by more advanced RNN-type architectures such as LSTMs and GRUs



Limitations of CNNs on Sequences

- ❑ Struggle to capture **long-range temporal dependencies**
- ❑ CNNs are **stateless** and don't carry information across time steps
 - ❑ Dilations (stride [t, t+2, t+4]) can help
- ❑ Unlike RNNs/LSTMs, they **don't** “remember” past context

