




Module 6

Recurrent Neural Networks



DL: RNNs




Sequences

- ❑ Enumerated collection of objects
- ❑ Repetitions in sequences are allowed and order matters.
- ❑ Time series, speech rely on changes over a time interval

$$\{a_n\} = \{a_1, a_2, a_3, \dots, a_n\}$$


Sequences Examples

- Stock price, financial instruments, inventory predictions



Sequences Examples


- ❑ Natural language processing, LLMs




Univariate vs Multivariate Sequences

- We may need to understand more than one sequence

**Univariate
Time Series**



**Multivariate
Time Series**



Feature Creation

- ❑ We can use a window size (width, length, step) and horizon
- ❑ This method is called the “sliding window”

```
# Window for one week with the target of predicting the next day (Bitcoin prices)
[123.654, 125.455, 108.584, 118.674, 121.338, 120.655, 121.795] -> [123.033]
[125.455, 108.584, 118.674, 121.338, 120.655, 121.795, 123.033] -> [124.049]
[108.584, 118.674, 121.338, 120.655, 121.795, 123.033, 124.049] -> [125.961]
```

window size = 7, horizon =
1

Feature Creation

- ❑ Using a window size (input width) and horizon (label width)
- ❑ This method is called the sliding window

```
time,  measure
1,     100
2,     110
3,     108
4,     115
5,     120
```

LAG 1 feature (window size = 1, horizon = 1)

```
X,      y
?,      100
100,    110
110,    108
108,    115
115,    120
120,    ?
```


Feature Creation

- ❑ Using a window size (input width) and horizon (label width)
- ❑ This method is called the sliding window

```
time,  measure
1,     100
2,     110
3,     108
4,     115
5,     120
```

LAG 1 feature (window size = 1, horizon = 2)

```
X1,  y1,  y2
?    100, 110
100, 110, 108
110, 108, 115
108, 115, 120
115, 120, ?
120, ?,   ?
```



Python

Creating Lags on Multiple Features

- ❑ Using the sliding window method on multiple predictors can help
- ❑ The windows used can be different

```
time,  measure
1,     100
2,     110
3,     108
4,     115
5,     120
```

LAG 1 feature (window size = 1, horizon = 2)


```
X1,  y1,  y2
?    100, 110
100, 110, 108
110, 108, 115
108, 115, 120
115, 120, ?
120, ?,   ?
```




Python

Recurrent Neural Networks

- ❑ RNN neurons contain a feedback loop to receive sequence inputs
- ❑ The neurons (units) can be rolled to account for a sequence.




$$y = f(w_{ih}x + b)$$



$$h_t = f(w_{ih}x + w_{hh}h_{t-1} + b)$$

Recurrent Neural Networks


- ❑ RNN neurons contain a feedback loop to receive sequence inputs
- ❑ The neurons (units) can be rolled to account for a sequence.



$$h_t = f(w_{ih}x + w_{hh}h_{t-1} + b)$$


Recurrent Neural Networks

- ❑ RNN neurons contain a feedback loop to receive sequence inputs
- ❑ The neurons (units) can be rolled to account for a sequence.




Recurrent Neural Networks

- Sequence to Sequence (seq2seq)




Recurrent Neural Networks

- Sequence to Vector (seq2vec)




Recurrent Neural Networks

□ Vector to Sequence (vec2seq)



Recurrent Neural Networks

- ❑ Elman network
- ❑ Tanh activation




(a) Simple RNN

HIDDEN STATE


Stacked RNN Architecture

- Stacked (deep) RNNs are composed of multiple RNNs stacked one above the other.



RNN Architecture

- ❑ We can construct layers of RNN neurons (units)
- ❑ We can also stack layers (need to return the sequences)




General Form of RNNs



Python

Bi-directional RNN Architecture

- ❑ Composed of two RNNs
- ❑ The input sequence (RNN # 1) and opposite direction (RNN # 2)






Python

CNN for Sequences

- ❑ CNNs excel at capturing **local dependencies** (e.g., trends, peaks).
- ❑ Effective in **repeating patterns** regardless of position in sequence.
- ❑ Fewer parameters than RNNs; allows **parallel computation**.



TimeSeriesGenerator

- ❑ Converts raw time series into supervised learning **format**
- ❑ Efficient for memory use and model training
- ❑ Supports **sliding window** logic

python

 Copy

```
TimeseriesGenerator(  
    data,                # Numpy array of time series  
    targets,             # Numpy array of target values  
    length=10,           # Number of time steps per input sample  
    sampling_rate=1,     # Period between samples  
    stride=1,            # Step between successive windows  
    batch_size=32,       # Number of samples per batch  
    shuffle=False        # Whether to shuffle samples  
)
```

TimeSeriesGenerator

- ❑ Converts raw time series into supervised learning **format**
- ❑ Efficient for memory use and model training
- ❑ Supports **sliding window** logic

Time	Stock	Feature 1	Feature 2	Feature ...	Feature N			
T+0		1	[some value]	[some value]	[some value]	[some value]		
T+1		1	[some value]	[some value]	[some value]	[some value]		
T+2		1	[some value]	[some value]	[some value]	[some value]		
T+3		1	[some value]	[some value]	[some value]	[some value]		
T+4		1	[some value]	[some value]	[some value]	[some value]	Sequence 1	
T+5		1	[some value]	[some value]	[some value]	[some value]		Sequence 2
T+0		2	[some value]	[some value]	[some value]	[some value]		Sequence 3
T+1		2	[some value]	[some value]	[some value]	[some value]		
T+2		2	[some value]	[some value]	[some value]	[some value]		
T+3		2	[some value]	[some value]	[some value]	[some value]		
T+4		2	[some value]	[some value]	[some value]	[some value]		
T+5		2	[some value]	[some value]	[some value]	[some value]		
T+0		3	[some value]	[some value]	[some value]	[some value]		
T+1		3	[some value]	[some value]	[some value]	[some value]		
T+2		3	[some value]	[some value]	[some value]	[some value]		

Review: Convolutions

- Suppose we have a sequence $I_{n \times p}$ and a kernel $K_{k \times l}$
- The resulting image $O_{n-k+1 \times p-l+1}$

7	2	3	3	8
4	5	3	8	4
3	3	2	8	4
2	8	7	2	7
5	4	4	5	4

*

1	0	-1
1	0	-1
1	0	-1

=

6		

$$\begin{aligned}
 &7 \times 1 + 4 \times 1 + 3 \times 1 + \\
 &2 \times 0 + 5 \times 0 + 3 \times 0 + \\
 &3 \times -1 + 3 \times -1 + 2 \times -1 \\
 &= 6
 \end{aligned}$$


$$I_{5 \times 5}$$

$$K_{3 \times 3}$$

$$O_{5-3+1 \times 5-3+1} = O_{3 \times 3}$$


Review: Convolutions

- ▣ Suppose we have a sequence $I_{n \times p}$ and a kernel $K_{k \times l}$
- ▣ The resulting image $O_{n-k+1} \times p-l+1$



Review: Pooling

- ❑ Pooling is also used to downsample the sequences.
- ❑ Pooling filters keep the important parts of the sequence.
- ❑ Max, Min and Average Pooling Filters are the most common





Python

Limitations of RNNs

- ❑ Struggle with long-term dependencies due to vanishing gradients.
 - ❑ Attention (weights or scores for each hidden state)
- ❑ Training can be slow and unstable.
- ❑ Can be replaced by more advanced RNN-type architectures such as LSTMs and GRUs



Limitations of CNNs on Sequences

- ❑ Struggle to capture **long-range temporal dependencies**
- ❑ CNNs are **stateless** and don't carry information across time steps
 - ❑ Dilations (stride $[t, t+2, t+4]$) can help
- ❑ Unlike RNNs/LSTMs, they **don't "remember" past context**

