




Module 6

Recurrent Neural Networks




DL: RNNs



LSTM Gates

- Gates allow LSTMs to control short- and long-term memories



$$\mathbf{F}_t = \sigma(\mathbf{X}_t \mathbf{W}_{xf} + \mathbf{H}_{t-1} \mathbf{W}_{hf} + \mathbf{b}_f)$$

$$\mathbf{I}_t = \sigma(\mathbf{X}_t \mathbf{W}_{xi} + \mathbf{H}_{t-1} \mathbf{W}_{hi} + \mathbf{b}_i)$$

$$\tilde{\mathbf{C}}_t = \tanh(\mathbf{X}_t \mathbf{W}_{xc} + \mathbf{H}_{t-1} \mathbf{W}_{hc} + \mathbf{b}_c)$$


$$\mathbf{C}_t = \mathbf{F}_t \odot \mathbf{C}_{t-1} + \mathbf{I}_t \odot \tilde{\mathbf{C}}_t$$

$$\mathbf{O}_t = \sigma(\mathbf{X}_t \mathbf{W}_{xo} + \mathbf{H}_{t-1} \mathbf{W}_{ho} + \mathbf{b}_o)$$

$$\mathbf{H}_t = \mathbf{O}_t \odot \tanh(\mathbf{C}_t).$$

Encoder-Decoder LSTM

- Seq2seq LSTM encoder-decoder networks




Attention Layers

- ❑ Traditional models like RNNs and LSTMs treat all time steps equally.
- ❑ In reality, some moments matter more than others.
- ❑ Goal: Learn to focus on important time steps dynamically.





What is Attention?

- ❑ A mechanism that lets models learn where to focus.
- ❑ Assigns a weight to each input element.
- ❑ Outputs a context vector summarizing important parts.



How Attention Works

- ❑ Score each time step with a small neural network.
- ❑ Apply softmax to get attention weights.
- ❑ Multiply each hidden state by its weight.
- ❑ Sum to get the context vector.



Attention Layers (Business Uses)

- ❑ Besides the regular use of LLMs
 - ❑ Forecasting
 - ❑ Past sales, promotions, holidays, weather
 - ❑ Attention focuses on:
 - ❑ Holiday seasons
 - ❑ Promotional weeks
 - ❑ Special weather events




Attention Layers Anatomy

- ❑ Example of Attention Layers:
 - ❑ Dense(1): Assigns a score to each time step.
 - ❑ Softmax: Normalizes into probabilities.
 - ❑ Reduce_sum: Dot Product creates the context vector.

```
attention_scores = Dense(1, activation='tanh')(lstm_out)
attention_weights = Activation('softmax')(attention_scores)
context_vector = tf.reduce_sum(attention_weights * lstm_out,
axis=1)
```

Benefits of Attention Layers

- ❑ Improved Accuracy: Focus on key time steps.
- ❑ Interpretability: Understand what the model "looks at."
- ❑ Scalability: Easier learning for long sequences.





Python

Multi-Head Attention Layers

- ❑ Introduced in Transformer models.
- ❑ Applies multiple attention heads in parallel.
- ❑ Each head learns to focus on different aspects of the input.
- ❑ Outputs from all heads are concatenated and linearly transformed.


Single Head	Multi-Head
One perspective on sequence	Many parallel perspectives
May miss complex patterns	Captures diverse relationships
Limited capacity	Higher representational power

Multi-Head Attention Layers

- ❑ Introduced in Transformer models.
- ❑ Applies multiple attention heads in parallel.
- ❑ Each head learns to focus on different aspects of the input.
- ❑ Outputs from all heads are concatenated and linearly transformed.


BERT

Encoder



GPT

Decoder





Python